



Detected indoor objects (MIT + NYU-Depth) vs the Amazon Home & Kitchen catalogue

Objects: the 1,068 nodes of the trained co-occurrence graph, verified equal to the default_fixed encoder checkpoint. Items: raw_graph.txt, one label per item, 1,815 distinct — the same file build_object_feat reads. Tiers are that function's own branch: exact name match, else nearest MiniLM node. The shipped corpus reaches 7,506 items exactly (51.8%) and 5,485 more by proximity (37.8%), leaving 1,512 (10.4%) unreachable. τ sets only the near/unreached split, never the exact share: 3.9% / 10.4% / 21.1% unreachable at $\tau = 0.5 / 0.6 / 0.7$.

(a) Catalogue reached by each scene corpus



(b) The catalogue's 22 largest labels are almost all exact object nodes



(c) Labels the vocabulary misses, and the proxy node each is assigned instead

