


MIT-Indoors vs NYU-Depth as object-graph corpora


MIT is 4.5x NYU's size and holds 83% of its vocabulary, yet every corpus lands within $0.53\text{e-}3$ of the others — inside this backbone's ~ 0.0017 resolution.

Raw MIT – NYU is $+0.53\text{e-}3$ ($t=+1.44$, n.s.); size-matched it falls to $+0.13\text{e-}3$ ($t=+0.62$). What little advantage MIT has is corpus size, not corpus quality — which is the arm the size-matched control exists to isolate. NYU also has the highest intrinsic encoder AUC and the lowest downstream score, the fourth instance in this study of intrinsic quality failing to predict recommendation.

(a) The corpora differ in scale



(b) ...and in vocabulary, but NYU is mostly inside MIT



(c) ...and not at all downstream. Tuned LightGCN, 3 seeds; no pair is resolvable

